

SAM Slide Suite — UI Walkthrough: Review & Rewrite

A stage-by-stage tour of the **Review** and **Rewrite** workflows, captured from the running app.

- **Document:** `AI-Powered Acquisition Document Processing Pipeline.pptx`
- **Model:** local Ollama (`kimi-k2.6:cloud`)
- **Review panel:** Skeptic-Proof Analyst · Executive Briefing · Legal/Compliance
- **Rewrite persona:** Executive Briefing, with **Advise** enabled


All screenshots are real runs (theme: *mixed* — dark chrome, light reading cards). The two workflows share one document and the same left config panel; only the persona-selection mode and the right output panel change per tab.

 Prefer a slideshow? Open ui-deck.html — a self-contained, light-blue HTML deck of these same stages with prev/next, keyboard arrows, and a clickable thumbnail rail.

Review

1. Empty state


Pick the Review tab, and the right panel explains the flow while you configure the left panel.



2. Configured — document + review panel


The deck is uploaded (extracted via `extract.py`) and confirmed in the Document section. In Review mode personas are **multi-select** (checkboxes) with Select All / Clear; here three are chosen. The Model picker

shows the detected Ollama model with a refresh (🔄) button.



3. Running — live per-persona progress

On Run Review, each persona streams status (queued → running → done) over SSE, with a progress bar and a live log. The log shows extraction, the auto-detected model context window, and each reviewer starting.



4. Results — synthesis + per-persona reports

When the panel finishes, a synthesis pass merges the reviews. The **Synthesis** tab (consensus findings, conflicts weighted by each persona's blind spots, unique catches, top-5 priority fixes) sits first, followed by one tab per persona. Export DOCX and Compare are below.

The screenshot displays the SAM Slide Suite interface. On the left, a sidebar contains sections for 'Document' (showing the file 'AI-Powered Acquisition Document Processing Pipeline.pptx'), 'Personas' (listing Sam, Commander Sam, Empathetic Holistic, Energetic Collaborative, Skeptic Proof Analyst, and Executive Briefing), and 'Model' (set to 'kimi-k2.6:cloud'). A 'Run Review' button is at the bottom of the sidebar. The main area shows the 'Review' tab selected, with sub-tabs for 'Skeptic Proof Analyst Persona', 'Executive Briefing Persona', and 'Legal Compliance Persona'. The 'Review Synthesis' section is active, displaying the title 'AI-Powered Acquisition Document Processing Pipeline.pptx — Review Synthesis' and the panel 'Panel: Skeptic Proof Analyst, Executive Briefing, Legal Compliance - 2026-06-17'. The 'Overall verdict' states: 'This presentation is currently a skeleton of titles and empty frames, lacking the substantive claims, metrics, and decision frames required for technical, executive, or compliance review. All three panelists concur that while the narrative sequence and federal acquisition domain scope provide a sound scaffold, the deck cannot be approved, funded, or audited until it is populated with quantified problem statements, substantiated AI evidence, deployable architecture, and a clear decision ask. The author must provide a text-extractable version and complete content before the panel can validate readiness.' The 'Consensus findings' section lists several issues: 'Extracted content is almost entirely empty; no reviewable substantive claims exist' (flagged by all three personas), 'Slide 3 ("SAM Problem") contains no concrete, quantified problem statement' (flagged by all three), 'Slide 6 ("SAM Deployment Architecture") lacks named components, boundaries, and data flows' (flagged by all three), 'Slide 7 ("Challenges / Questions") is empty and omits risks, dependencies, and the decision ask' (flagged by all three), and 'Slide 1 AI claim is unsubstantiated and lacks a basis or scope boundary'.

5. Compare — two reports side by side

The Compare overlay puts any two reports next to each other (each pane scrolls independently) for quick cross-persona reading.

Compare:

Synthesis

 vs

Skeptic Proof Analyst Persona

Review complete

AI-Powered Acquisition Document Processing Pipeline.pptx — Review Synthesis

Panel: Skeptic Proof Analyst, Executive Briefing, Legal Compliance · 2026-06-17

Overall verdict

This presentation is currently a skeleton of titles and empty frames, lacking the substantive claims, metrics, and decision frames required for technical, executive, or compliance review. All three panelists concur that while the narrative sequence and federal acquisition domain scope provide a sound scaffold, the deck cannot be approved, funded, or audited until it is populated with quantified problem statements, substantiated AI evidence, deployable architecture, and a clear decision ask. The author must provide a text-extractable version and complete content before the panel can validate readiness.

Consensus findings

Extracted content is almost entirely empty; no reviewable substantive claims exist

- Flagged by: Skeptic Proof Analyst, Executive Briefing, Legal Compliance
- Fix: Populate every slide with full text, bullets, diagram annotations, and speaker notes; provide a text-extractable transcript or structured transcript of all content so that technical, executive, and compliance reviewers can audit claims and controls.

Slide 3 ("SAM Problem") contains no concrete, quantified problem statement

- Flagged by: Skeptic Proof Analyst, Executive Briefing, Legal Compliance
- Fix: Rewrite using an inverted-pyramid headline followed by baseline metrics: current processing time, error rate, monthly volume, cost, measurement method, and time period. Example: "Current manual processing averages 4.2 hours per acquisition file, with a 12% routing error rate measured over Q1-Q2 2025."

Slide 6 ("SAM Deployment Architecture") lacks named components, boundaries, and data flows

- Flagged by: Skeptic Proof Analyst, Executive Briefing, Legal Compliance
- Fix: Add named components, responsibilities, network boundaries, directional data-flow arrows, and explicit controls where sensitive acquisition data crosses boundaries; frame the slide around a decision input (e.g., "Approve three-layer architecture to eliminate single points of failure").

Slide 7 ("Challenges / Questions") is empty and omits risks, dependencies, and the decision ask

Flagged by Skeptic Proof Analyst, Executive Briefing, Legal Compliance

SAM-Flow AI-Powered Acquisition Document Processing Pipeline — Review by Skeptic Proof Analyst

The extracted content is almost entirely slide titles and section headers. No problem statement, no evidence, no architecture description, and no figures appear in the text provided. I can only review what is visible. What is visible is a title that asserts an AI capability without proof, followed by a series of placeholder headings. Until the slides contain actual claims, numbers, and boundaries, this deck is not ready for technical review.

Strengths

- The slide sequence follows a defensible narrative order: problem, process flow, architecture, then challenges.
- The scope is anchored to a specific business domain—acquisition document processing—rather than a generic AI proof-of-concept.

Findings

1. Unsubstantiated AI capability claim — Slide 1

Severity: Critical

Issue: The title "AI-Powered Acquisition Document Processing Pipeline" asserts that artificial intelligence is central to the solution, yet the extraction contains no model name, service provider, inference architecture, or performance benchmark. The claim is presented as fact without evidence or an explicit assumption flag.

Recommendation: Add a subtitle or adjacent text naming the specific AI component (e.g., a named LLM, classifier, or OCR engine) and the primary capability it provides. If the AI element is not yet selected, flag it as an assumption: "AI component assumed; evaluation in progress."

2. Load-bearing problem statement is missing — Slide 3

Severity: Major

Issue: The slide title "SAM Problem" suggests the problem definition lives here, but no problem description appears in the extraction. Without the problem's scope, volume, error rate, or cost, the audience cannot evaluate whether the proposed pipeline is proportionate or effective.


Recommendation: State the problem in concrete terms. Include current processing time, error rate, volume growth, or compliance gap. Include the measurement method and the time period. Example: "Current manual processing averages 4.2 hours per acquisition file, with a 12% routing error rate measured over Q1-Q2 2025."

3. Architecture slide contains no architecture — Slide 6

Rewrite


6. Configured — single persona + Advise

Switching to Rewrite keeps the same document. Personas become **single-select** (radio). The Rewrite-only  **Advise** toggle is enabled here, so the persona will draft proposed content for each gap the review flagged.



7. Running

Run Rewrite streams progress while the persona rewrites the whole deck. It automatically folds in the prior review's findings for this document.



8. Results — rewritten slides with NEEDS / DRAFT

Each slide becomes a tab showing the rewritten on-slide content plus speaker notes. Gaps the review surfaced appear as amber **NEEDS** markers; because Advise is on, each is followed by a purple **DRAFT** proposal — a labeled, unverified starting point to confirm or replace. Export PPTX (real speaker-notes pages) and Export DOCX are below.

The screenshot displays the SAM Slide Suite interface. On the left, a sidebar contains sections for Document, Personas, Advise, and Model. The main area shows a slide titled "Rewrite by: Executive Briefing Persona" with tabs for Slide 2 through Slide 7. The slide content includes a list of bullet points and a paragraph, both containing NEEDS and DRAFT markers. Below the slide content, there is a section for "Speaker Notes" and a large "Placeholders (50)" section. At the bottom, there are buttons for "Export PPTX" and "Export DOCX".

Document

- ✓ AI-Powered Acquisition Document Processing Pipeline.pptx

Personas

- Sam: Narrator and guide — the voice of the architect...
- Commander Sam: The voice that refuses to waste your time or let yo...
- Empathetic Holistic: Compassionate systems guide — a voice that ho...
- Energetic Collaborative: Teammate and guide — the voice of the architect...
- Skeptic Proof Analyst: Evidence-first analyst — the voice that already as...
- Executive Briefing: Time compressed decision maker — the voice that...

Advise

- ✓ Draft content for [NEEDS:] gaps

The persona proposes [DRAFT:] values for each gap the review flagged — clearly marked for you to confirm.

Model

kimi-k2.6:cloud

Context (optional)

Rewrite by: Executive Briefing Persona

Slide 2 Slide 3 Slide 4 Slide 5 Slide 6 Slide 7

- Enterprise Architecture Division
- Scope: [NEEDS: confirm SAM refers to System for Award Management and define agency scope] [DRAFT: System for Award Management structured document intake]
- AI component: [NEEDS: name specific LLM, classifier, or OCR engine and evidence] [DRAFT: LLM-based classifier assumed; vendor evaluation in progress]

--- Speaker Notes ---

This is the SAM-Flow pilot proposal. Scope is [NEEDS: confirm SAM refers to System for Award Management and define agency scope] [DRAFT: System for Award Management structured-document intake]. The AI component is [NEEDS: name specific LLM, classifier, or OCR engine and evidence] [DRAFT: an assumed large language model classifier; vendor evaluation is in progress]. We are not presenting a selected vendor as fact. Everything that follows justifies the pilot and defines the guardrails.

Placeholders (50): confirm SAM refers to System for Award Management and define agency scope · name specific LLM, classifier, or OCR engine and evidence · confirm SAM refers to System for Award Management and define agency scope · name specific LLM, classifier, or OCR engine and evidence · pilot duration and budget ceiling · cost-of-delay metric per month · decision deadline · current processing time per file with period · routing error rate with method and period · monthly volume figure · monthly or per-file cost with period · current processing time per file with period · routing error rate with method and period · monthly volume figure · monthly or per-file cost with period · current SAM intake process steps and handoff count · target automated process steps · handoffs to eliminate · current SAM intake process steps and handoff count · target automated process steps · handoffs to eliminate · ingestion component name and responsibility · inference component name, AI model boundary, and hosting posture · output component name and responsibility · network boundary crossing descriptions and enforcement controls · documented non-AI deterministic alternative · ingestion component name and responsibility · inference component name, AI model boundary, and hosting posture · output component name and responsibility · network boundary crossing descriptions and enforcement controls · documented non-AI deterministic alternative · explicit ask wording and scope · quantified expected benefit · cost-of-delay metric · risk mitigation detail · dependency list · open question, owner, due date · open question, owner, due date · open question, owner, due date · data residency and hosting posture · applicable control framework citations · 48 drafted

Export PPTX **Export DOCX**

Rewrite Deck

9. Another slide

Stepping through the slide tabs shows the same structure per slide — rewritten content, speaker notes, and any NEEDS/DRAFT markers.

The screenshot displays the SAM Slide Suite application interface. The top navigation bar includes 'Review', 'Rewrite' (selected), 'Narrate', and 'Produce'. A 'Rewrite complete' button is visible in the top right. The left sidebar contains a 'Document' section with the file 'AI-Powered Acquisition Document Processing Pipeline.pptx', a 'Personas' section with roles like 'Sam', 'Commander Sam', 'Empathetic Holistic', 'Energetic Collaborative', 'Skeptical Proof Analyst', and 'Executive Briefing', and an 'Advise' section with a checkbox for 'Draft content for [NEEDS:] gaps'. The main area shows the 'Rewrite by: Executive Briefing Persona' for 'Slide 3'. The slide content includes a list of bullet points with placeholders for 'NEEDS' and 'DRAFT' text, followed by 'Speaker Notes' and a detailed paragraph. At the bottom of the slide area are 'Export PPTX' and 'Export DOCX' buttons. A 'Rewrite Deck' button is located at the bottom left of the interface.

How these were captured

Headless Chrome (via `puppeteer-core`) drove the built app served by Express on `:3001`, uploading the deck and walking each stage against a local Ollama model. Reviews ran sequentially (`REVIEW_CONCURRENCY=1`) for deterministic capture. To regenerate, serve the app, then drive the same steps with any browser-automation tool.